

# SIMPLE LINEAR REGRESSION

## 1. SELECTION OF DATASET:

The dataset used here is 'Glass Identification Database'. This dataset is used to predict the glass type from chemical properties.

Dimension: 214 rows and 10 columns

```
> library(mlbench)
> data(Glass)
> head(Glass) #To see the first six rows of the dataset
```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
1	1.52101	13.64	4.49	1.10	71.78	0.06	8.75	0	0.00	1
2	1.51761	13.89	3.60	1.36	72.73	0.48	7.83	0	0.00	1
3	1.51618	13.53	3.55	1.54	72.99	0.39	7.78	0	0.00	1
4	1.51766	13.21	3.69	1.29	72.61	0.57	8.22	0	0.00	1
5	1.51742	13.27	3.62	1.24	73.08	0.55	8.07	0	0.00	1
6	1.51596	12.79	3.61	1.62	72.97	0.64	8.07	0	0.26	1

```
> summary(Glass) #To see the summary of the dataset
```

	RI	Na	Mg	Al	Si	K	Ca
Min.	:1.511	Min. :10.73	Min. :0.000	Min. :0.290	Min. :69.81	Min. :0.0000	Min. : 5.430
1st Qu.:	:1.517	1st Qu.:12.91	1st Qu.:2.115	1st Qu.:1.190	1st Qu.:72.28	1st Qu.:0.1225	1st Qu.: 8.240
Median :	:1.518	Median :13.30	Median :3.480	Median :1.360	Median :72.79	Median :0.5550	Median : 8.600
Mean :	:1.518	Mean :13.41	Mean :2.685	Mean :1.445	Mean :72.65	Mean :0.4971	Mean : 8.957
3rd Qu.:	:1.519	3rd Qu.:13.82	3rd Qu.:3.600	3rd Qu.:1.630	3rd Qu.:73.09	3rd Qu.:0.6100	3rd Qu.: 9.172
Max. :	:1.534	Max. :17.38	Max. :4.490	Max. :3.500	Max. :75.41	Max. :6.2100	Max. :16.190

	Ba	Fe	Type
Min.	:0.000	Min. :0.00000	1:70
1st Qu.:	:0.000	1st Qu.:0.00000	2:76
Median :	:0.000	Median :0.00000	3:17
Mean :	:0.175	Mean :0.05701	5:13
3rd Qu.:	:0.000	3rd Qu.:0.10000	6: 9
Max. :	:3.150	Max. :0.51000	7:29

```
> str(Glass) #To see the structure of the dataset
```

```
'data.frame': 214 obs. of 10 variables:
 $ RI : num 1.52 1.52 1.52 1.52 1.52 ...
 $ Na : num 13.6 13.9 13.5 13.2 13.3 ...
 $ Mg : num 4.49 3.6 3.55 3.69 3.62 3.61 3.6 3.61 3.58 3.6 ...
 $ Al : num 1.1 1.36 1.54 1.29 1.24 1.62 1.14 1.05 1.37 1.36 ...
 $ Si : num 71.8 72.7 73 72.6 73.1 ...
 $ K : num 0.06 0.48 0.39 0.57 0.55 0.64 0.58 0.57 0.56 0.57 ...
 $ Ca : num 8.75 7.83 7.78 8.22 8.07 8.07 8.17 8.24 8.3 8.4 ...
 $ Ba : num 0 0 0 0 0 0 0 0 0 0 ...
 $ Fe : num 0 0 0 0 0 0.26 0 0 0 0.11 ...
 $ Type: Factor w/ 6 levels "1","2","3","5",...: 1 1 1 1 1 1 1 1 1 1 ...
```

The structure of the data set gives a detailed description of the data types and the levels for the variables which have factors.

```
> any(is.na(Glass)) #To check whether there are any null values in the dataset
[1] FALSE
```

## 2. CORRELATION ANALYSIS OF THE DATASET:


```
> #The following libraries are used for Exploratory Data Analysis
> library(ggplot2)
> library(ggthemes)
> library(dplyr)
```

```
> #Grabbing only the numeric columns as we can't see correlation between the categorical variables
> num.cols<- sapply(Glass,is.numeric)
> #Filtering to numeric columns for correlation
> cor.data<-cor(Glass[, num.cols])
> cor.data
```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe
RI	1.0000000000	-0.19188538	-0.122274039	-0.40732603	-0.54205220	-0.289832711	0.8104027	-0.0003860189	0.143009609
Na	-0.1918853790	1.000000000	-0.273731961	0.15679367	-0.06980881	-0.266086504	-0.2754425	0.3266028795	-0.241346411
Mg	-0.1222740393	-0.27373196	1.000000000	-0.48179851	-0.16592672	0.005395667	-0.4437500	-0.4922621178	0.083059529
Al	-0.4073260341	0.15679367	-0.481798509	1.000000000	-0.00552372	0.325958446	-0.2595920	0.4794039017	-0.074402151
Si	-0.5420521997	-0.06980881	-0.165926723	-0.00552372	1.000000000	-0.193330854	-0.2087322	-0.1021513105	-0.094200731
K	-0.2898327111	-0.26608650	0.005395667	0.32595845	-0.19333085	1.000000000	-0.3178362	-0.0426180594	-0.007719049
Ca	0.8104026963	-0.27544249	-0.443750026	-0.25959201	-0.20873215	-0.317836155	1.000000000	-0.1128409671	0.124968219
Ba	-0.0003860189	0.32660288	-0.492262118	0.47940390	-0.10215131	-0.042618059	-0.1128410	1.000000000	-0.058691755
Fe	0.1430096093	-0.24134641	0.083059529	-0.07440215	-0.09420073	-0.007719049	0.1249682	-0.0586917554	1.000000000

```
> #For proper data visualization we use the 'corrgram' package and the 'corrplot' package
> library(corrgram)
> library(corrplot)
```

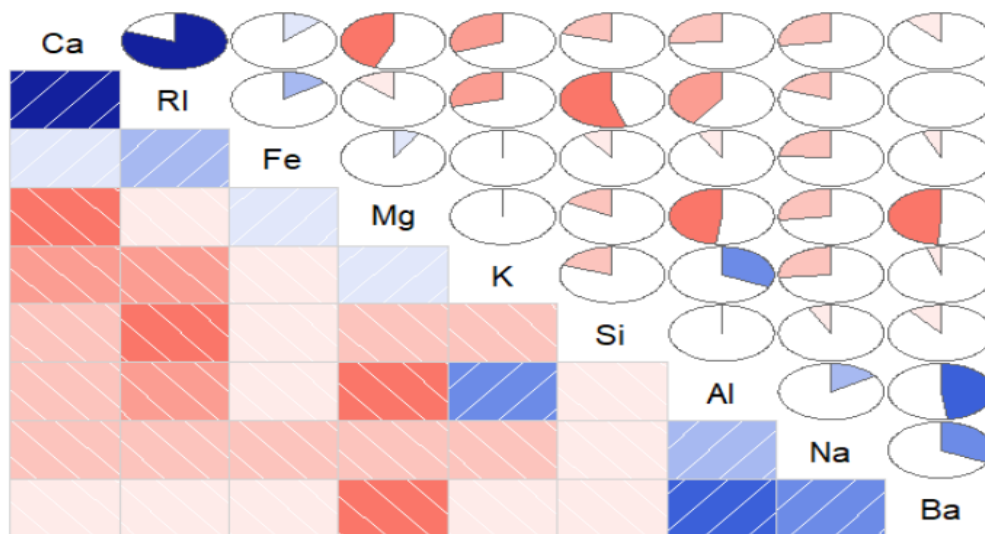
```
> #Now we perform the correlation plot and see what we can infer from that
> corrplot(cor.data, method='color')
```



We can see from the above correlation plot that white indicates zero correlation, the shades from white to blue indicates positive correlation and from white to red indicates negative correlation. Positive correlation indicates that as the value of one variable increases, the value


of the other variable also increases, i.e., directly proportional, while Negative correlation indicates that as the value of one variable increases the value of the other variable will decrease, i.e., indirectly proportional. We can see that Ca and RI are highly correlated to each other. This means that higher the refractive index(RI), higher is the amount of Calcium(Ca). We are going to predict the glass type with Ca based on RI. Here Ca is a dependent variable and RI is an independent variable.

```
> #Now we perform the correlogram and see what we can infer from that
> corrgram(Glass ,order=TRUE, lower.panel=panel.shade, upper.panel=panel.pie, text.panel=panel.txt)
```



In R, correlograms are implemented using the ‘corrgram’ function. It shows the graph of the correlation matrix and is very useful for highlighting the most correlated values. The results of this plot can be interpreted in the same way as ‘corrplot’. Here blue is positive correlation and pink is negative correlation.

```
> #Histogram of the amount of Ca
> ggplot(Glass,aes(x=Ca))+ geom_histogram(bins=20, alpha=0.5, fill='pink') + theme_minimal()
> |
```



As we have to predict the amount of Ca so we draw a histogram of it. From the graph we can see that the high amount of Ca is below the mean of the Ca and there are much more values or observations beyond the mean value. So we can say that this graph is positively skewed.

### 3. SIMPLE LINEAR REGRESSION:

```
> #We need to split our data into a training set and a testing set in order to test our accuracy, so we can do this using the caTools library
> library(caTools)
>
> #Splitting up the sample for training and testing and assigns boolean values to a new column
> sample <- sample.split(Glass$Ca, SplitRatio = 0.70)
>
> #Training data
> train_data = subset(Glass, sample == TRUE)
>
> #Testing data
> test_data= subset(Glass, sample == FALSE)
>
> #Training the model
> model <- lm(Ca ~ RI, train_data)
>
> summary(model)
```

Call:  
lm(formula = Ca ~ RI, data = train\_data)

Residuals:

Min	1Q	Median	3Q	Max
-3.1264	-0.4039	-0.0320	0.3379	2.3436

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-619.77	33.32	-18.60	<2e-16 ***
RI	414.05	21.94	18.87	<2e-16 ***

---  
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.7743 on 147 degrees of freedom  
Multiple R-squared: 0.7078, Adjusted R-squared: 0.7058  
F-statistic: 356.1 on 1 and 147 DF, p-value: < 2.2e-16

The residuals are the difference between the actual values of the variable we are predicting and predicted values from our regression.

The stars are for significance levels, with the number of asterisks displayed according to the p-value computed. \*\*\* indicates high significance. In this case, \*\*\* indicates that there is high significance between RI and Ca.

The estimated coefficient is the value of slope calculated by the regression.

Standard Error of the Coefficient Estimate is the measure of the variability in the estimate for the coefficient.

t-value shows how many standard deviations the coefficient is far away from zero. Further it is away from zero, stronger the relationship between the variables.


p-value shows whether the overall model is significant or not. Coefficients having p-value less than the level of significance are said to be statistically significant. In this case we can say that the variable RI is statistically significant for the prediction of amount of Ca.

R-squared is an overall measure of the strength of association and adjusted R-squared gives a more proper estimate of the R-squared value for the population. Its value shows that 70.58 % of the variance in the amount of Ca can be predicted from RI.

```

> #Visualizing the model
>
> res<- residuals(model)      #Grabbing the residuals
> res<- as.data.frame(res)    #Converting the residuals to dataframe for ggplot
> head(res)
      res
1 -1.2553089
3 -0.2254548
4 -0.3982465
6  0.1556359
7 -0.3530154
10 -0.1727012
>
> #Histogram of residuals
> ggplot(res, aes(res)) + geom_histogram(fill='pink', alpha=0.5)
`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

```



We have to remove the negative values from here.

```

> #Predictions
> Ca.predictions <- predict(model,test_data)
> results<- cbind(Ca.predictions,test_data$Ca)
> colnames(results) <- c('pred','real')
> results<- as.data.frame(results)
> results

```

	pred	real
2	8.597544	7.83
5	8.518875	8.07
8	8.576842	8.24
9	9.247600	8.30
11	7.810852	8.09
23	8.494032	8.70
30	8.692775	8.49
33	8.655511	8.56
43	8.672073	8.59
44	10.456622	9.74
46	9.173071	9.00
48	12.348823	9.82
57	6.336839	8.35
62	9.491889	8.67
68	10.216474	9.85
69	10.216474	9.82
71	7.823273	7.36
73	7.901943	7.83
79	7.984752	7.94
83	8.121388	8.10
84	7.906083	8.05
86	8.034438	8.21
87	7.802571	8.03
90	8.096545	8.08
91	8.928783	8.96
97	8.928783	9.13
98	8.523015	8.90
102	8.469189	9.23

106	11.553850	13.24
107	14.245165	13.30
109	10.506308	11.52
110	8.833552	10.99
120	8.146231	7.96
121	8.941204	8.42
127	8.208339	8.60
129	9.868673	9.57
142	8.970188	8.41
143	8.187636	8.23
145	8.179355	8.81
148	7.972331	8.33
150	8.108967	8.53
158	10.088119	9.65
160	8.742461	8.81
161	8.891518	8.99
163	10.460762	9.14
164	7.574844	5.87
166	10.295143	11.41
168	9.458765	11.53
169	8.204198	10.17
173	6.775731	6.93
174	9.765161	12.50
175	9.827268	9.70
177	9.193774	9.57
182	9.123386	9.95
185	5.922791	6.65
186	5.989039	5.43
190	11.098397	8.61
191	7.984752	8.67
192	7.939207	8.76
196	7.703199	9.07
197	7.748745	9.41
201	7.550001	8.34
203	7.574844	8.39
204	8.171074	8.28
206	8.477470	8.61

```
> #To remove negative predictions and replace it with 0
> to_zero <- function(x){
+   if(x <0) {
+     return(0)
+   }else{
+     return(x)
+   }
+ }
>
> results$pred <- sapply(results$pred,to_zero)
> results
```

	pred	real
2	8.597544	7.83
5	8.518875	8.07
8	8.576842	8.24
9	9.247600	8.30
11	7.810852	8.09
23	8.494032	8.70
30	8.692775	8.49
33	8.655511	8.56
43	8.672073	8.59
44	10.456622	9.74
46	9.173071	9.00
48	12.348823	9.82
57	6.336839	8.35
62	9.491889	8.67
68	10.216474	9.85
69	10.216474	9.82
71	7.823273	7.36
73	7.901943	7.83
79	7.984752	7.94
83	8.121388	8.10
84	7.906083	8.05
86	8.034438	8.21
87	7.802571	8.03
90	8.096545	8.08
91	8.928783	8.96

97	8.928783	9.13
98	8.523015	8.90
102	8.469189	9.23
106	11.553850	13.24
107	14.245165	13.30
109	10.506308	11.52
110	8.833552	10.99
120	8.146231	7.96
121	8.941204	8.42
127	8.208339	8.60
129	9.868673	9.57
142	8.970188	8.41
143	8.187636	8.23
145	8.179355	8.81
148	7.972331	8.33
150	8.108967	8.53
158	10.088119	9.65
160	8.742461	8.81
161	8.891518	8.99
163	10.460762	9.14
164	7.574844	5.87
166	10.295143	11.41
168	9.458765	11.53
169	8.204198	10.17
173	6.775731	6.93
174	9.765161	12.50
175	9.827268	9.70
177	9.193774	9.57
182	9.123386	9.95
185	5.922791	6.65
186	5.989039	5.43
190	11.098397	8.61
191	7.984752	8.67
192	7.939207	8.76
196	7.703199	9.07
197	7.748745	9.41
201	7.550001	8.34
203	7.574844	8.39
204	8.171074	8.28
206	8.477470	8.61

```
> #Evaluating the prediction values by the method of MSE(Mean Squared Error)
> mse <- mean((results$real - results$pred)^2)
> print(mse)
[1] 0.9650583
>
> #Evaluating the prediction values by the method of RMSE(Root Mean Squared Error)
> mse^0.5
[1] 0.9823738
>
> #Or we can just use the R-Squared Value for the model which gives the accuracy of the model
> SSE <- sum((results$pred - results$real)^2) #Sum Square of Errors
> TSS <- sum((mean(Glass$Ca) - results$real)^2) #Total Sum of Squares
> R2 <- 1-SSE/TSS #R-Squared
> R2
[1] 0.5165802
```

R-Squared (Coefficient of Determination) - This value lies between 0 and 1, and the higher it is, the better the model fits the data set.

Our main aim is to find those variables who give the lowest RMSE value and the highest R-Squared value.

The R-Squared for the training set is 51.65%. It means that the model can explain more than 51.65% of the variation.